Statistical Manifold Geometry of Regime Dynamics:

A Universal Attractor Framework for Changepoint Detection, Trust Quantification, and Decision-Making Under Non-Stationarity

Samyak Jain

NIT Rourkela

jsamyak2068@gmail.com

March 2026 — Preprint

Abstract

We introduce the **Statistical Manifold Attractor (SMA) Framework**, a geometric approach to regime dynamics in non-stationary time series. Rather than treating regimes as discrete latent labels, we represent the statistical state of a process at time t as a vector (t) of rolling statistical properties residing on a smooth parameter manifold $\Theta \subseteq \mathbb{R}^k$. The velocity $\boldsymbol{v}(t) = (t) - (t{-}1)$, speed $s(t) = \|\boldsymbol{v}(t)\|$, and net displacement $\boldsymbol{d}(t, W) = (t) - (t{-}W)$ define the local kinematics of regime change in continuous, coordinate-free terms.

We establish three main results. **(i) Speed as a transition detector:** $s(t)$ constitutes a real-time, model-free regime-transition signal that cleanly identifies all known crisis events across four independent domains without parametric assumptions. **(ii) Universal mean reversion:** the displacement persistence score $P(t, W) = \langle \hat{\boldsymbol{d}}(t, W), \hat{\boldsymbol{d}}(t{+}W, W) \rangle$ satisfies $\mathbb{E}[P(t, W)] < 0$ with high significance ($p < 0.001$) in all four domains tested — finance, solar physics, epidemiology, and climatology — confirming that the statistical parameter trajectory is universally mean-reverting toward a stable attractor $\boldsymbol{\theta}^*$. We prove this is a necessary consequence of the boundedness of the system. **(iii) Data-trust score:** the attractor distance $D(t) = \|(t) - \boldsymbol{\theta}^*\|$ provides a principled basis for a continuous data-trust score $\alpha(t) \in [0, 1]$ that quantifies in real time how much weight to assign to current observations versus structural priors.

All results are cross-domain and parameter-free.

Contents

1 Introduction **4**

2 Mathematical Framework **5**

2.1 The Observation Process 5

2.2 The Statistical Feature Map and Parameter Manifold 5

2.3 Trajectory and Kinematics 6

2.4 Regimes and Regime Boundaries 7

2.5 The Statistical Attractor 7

2.6 Displacement Persistence 8

2.7 PCA Reduction and Dimensionality Collapse 8

3 Theoretical Results **9**

3.1 Assumptions 9

3.2 Speed as a Regime Transition Signal 9

3.3 The Universal Attractor Theorem 10

3.4 Ornstein-Uhlenbeck Representation 11

3.5 Attractor Escape and Crisis Regimes 12

3.6 The Dimensionality Collapse Theorem 12

3.7 Attractor Migration 13

4 The Data-Trust Score $\alpha(t)$ **13**

4.1 Motivation 13

4.2 Formal Definition 13

4.3 Blended Prediction Under Uncertainty 14

4.4 Calibration 14

5 Empirical Methodology **14**

5.1 Datasets 14

5.2 Statistical Tests 14

6 Empirical Results **15**

6.1 Cross-Domain Displacement Persistence 15

6.2 Escape Phase Detection: Bear 2020 Q1 15

6.3 Dimensionality Collapse 15

6.4 Attractor Migration in Climate Data 16

7 Discussion **16**

7.1 What is New 16

7.2 Relation to Existing Work 16

7.3 Limitations 16

8 Conclusion **17**

1 Introduction

A central challenge in applied statistics, machine learning, and decision science is *acting on data whose generating distribution is changing.* In finance this is market regime shift; in epidemiology it is epidemic wave transition; in climate science it is tipping-point crossing; in manufacturing it is equipment degradation. Every domain faces the same structural problem: the statistical model was calibrated on historical data, but the current process has departed from that history.

Existing approaches address this through two paradigms. *Changepoint detection* (PELT [Killick et al., 2012], BOCPD [Adams & MacKay, 2007], gfpop [Runge et al., 2020]) identifies the discrete times at which the data-generating process changed. *Regime-switching models* [Hamilton, 1989] represent the process as switching among latent discrete states. Both paradigms provide a *classification* answer: which regime are we in, and when did we switch?

We propose a complementary formulation. Instead of classifying regime membership, we ask:

Where in statistical parameter space is the system, how fast is it moving, and toward what does it return?

This reframing shifts the primary object of study from *labels* to *trajectories*, and from *classification* to *geometry.* The statistical parameter vector (t) traces a path on a manifold Θ, and the geometric properties of that path — its speed, curvature, distance from equilibrium, and persistence of direction — constitute a richer and more actionable description of non-stationarity than any discrete label.

Our three principal contributions are:

1. A rigorous mathematical framework defining the parameter manifold, trajectory, kinematics, attractor, and regime regions (§2).
2. Three formal theoretical claims proved under explicit assumptions, including the Universal Attractor Theorem (§3).
3. Empirical validation across four independent domains with a cross-domain comparison table and the derivation of the data-trust score $\alpha(t)$ (§6–§4).

2 Mathematical Framework

2.1 The Observation Process

Definition 2.1 (Observation process)**.** Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space. An *observation process* is a measurable stochastic process

$$X : \{1, 2, \ldots, T\} \times \Omega \longrightarrow \mathbb{R}^d, \qquad d \geq 1.$$

We write X_t for the observation at time t. The process is *locally stationary* on an interval $[s, e]$ if the joint distribution of $(X_s, \ldots, X_e)$ can be approximated by a stationary process with slowly varying parameters.

Definition 2.2 (Observation window)**.** For width $W \in \mathbb{N}$ and time $t > W$, the *observation window* ending at t is the finite sequence

$$\mathbf{X}_{t:W} \;=\; (X_{t-W}, X_{t-W+1}, \ldots, X_{t-1}) \;\in\; \mathbb{R}^{d \times W}.$$

2.2 The Statistical Feature Map and Parameter Manifold

Definition 2.3 (Statistical feature map)**.** A *statistical feature map* of order k is a measurable function

$$\Phi : \mathbb{R}^{d \times W} \longrightarrow \mathbb{R}^k$$

that extracts k statistical properties from an observation window.

In our implementation with d=1 and k=12, we use:

$$\Phi\big(\mathbf{X}_{t:W}\big) \;=\; \big[\mu_t,\ \sigma_t^2,\ \gamma_t,\ \kappa_t,\ \rho_t^{(1)},\ \rho_t^{(2)},\ \rho_t^{(5)},\ \nu_t,\ \sigma_{\downarrow,t}^2,\ \mathrm{MAD}_t,\ \mathrm{DD}_t,\ G_t\big]^\top \tag{1}$$

where the components are defined as follows. Let $r_i = X_i$ denote the observations in the window.

$$\mu_t = \frac{1}{W}\sum_i r_i \qquad \sigma_t^2 = \frac{1}{W-1}\sum_i (r_i - \mu_t)^2$$

$$\gamma_t = \frac{(1/W)\sum_i (r_i - \mu_t)^3}{\sigma_t^3} \qquad \kappa_t = \frac{(1/W)\sum_i (r_i - \mu_t)^4}{\sigma_t^4} - 3$$

$$\rho_t^{(\ell)} = \mathrm{Corr}(r_i, r_{i+\ell}) \qquad \nu_t = \mathrm{Std}\big(|r_i|\big)$$

$$\sigma_{\downarrow,t}^2 = \frac{1}{n^-}\sum_{r_i<0} (r_i - \mu_t)^2 \qquad \mathrm{MAD}_t = \frac{1}{W}\sum_i |r_i - \mu_t|$$

$$\mathrm{DD}_t = \min_i \frac{\prod_{j\leq i}(1+r_j) - \max_{j'\leq i}\prod_{j\leq j'}(1+r_j)}{\max_{j'\leq i}\prod_{j\leq j'}(1+r_j)} \qquad G_t = \frac{\mathbb{E}[r_i \mid r_i > 0]}{|\mathbb{E}[r_i \mid r_i < 0]|}$$

Definition 2.4 (Statistical feature vector)**.** The *statistical feature vector* at time t is

$$(t) \;=\; \Phi\big(\mathbf{X}_{t:W}\big) \;\in\; \mathbb{R}^k.$$

Definition 2.5 (Statistical parameter manifold)**.** The *statistical parameter manifold* is the image of the feature map:

$$\Theta \;=\; \big\{ \in \mathbb{R}^k := \Phi(\mathbf{x}) \text{ for some } \mathbf{x} \in \mathbb{R}^{d\times W} \big\} \;\subseteq\; \mathbb{R}^k.$$

Under mild regularity conditions on Φ, Θ is a smooth submanifold of $\mathbb{R}^k$.

Remark 2.1 (Geometry of Θ)*.* The natural metric on the space of probability distributions is the Fisher–Rao metric [Amari, 1985], whose local form is given by the Fisher information matrix:

$$\mathcal{I}_{ij}() \;=\; \mathbb{E}\left[\frac{\partial \log p_{(X)}}{\partial \theta_i}\frac{\partial \log p_{(X)}}{\partial \theta_j}\right].$$

All distances stated in this paper use the Euclidean norm in $\mathbb{R}^k$ on the standardised feature vector. The Fisher–Rao metric is a natural refinement left for future work.

2.3 Trajectory and Kinematics

Definition 2.6 (Trajectory)**.** The *statistical trajectory* of the process up to time T is

$$\tau \;=\; \big((1),(2),\ldots,(T)\big) \;\in\; \Theta^T.$$

Definition 2.7 (Velocity and speed)**.** The *velocity vector* at time $t \geq 2$ is

$$\boldsymbol{v}(t) \;=\; (t) - (t-1) \;\in\; \mathbb{R}^k.$$

The *instantaneous speed* is

$$s(t) \;=\; \|\boldsymbol{v}(t)\|_2 \;=\; \sqrt{\sum_{i=1}^{k}\big[\theta_i(t) - \theta_i(t-1)\big]^2}.$$

Definition 2.8 (Acceleration)**.** The *acceleration vector* at time $t \geq 3$ is

$$\boldsymbol{a}(t) \;=\; \boldsymbol{v}(t) - \boldsymbol{v}(t-1) \;=\; (t) - 2(t-1) + (t-2).$$

The *acceleration magnitude* is $a(t) = \|\boldsymbol{a}(t)\|_2$.

Definition 2.9 (Displacement and direction)**.** For window size $W \geq 1$ and $t > W$, the

net displacement vector is

$$\boldsymbol{d}(t, W) \;=\; (t) - (t - W) \;\in\; \mathbb{R}^k.$$

When $\|\boldsymbol{d}(t, W)\|_2 > 0$, the *unit displacement direction* is

$$\hat{\boldsymbol{d}}(t, W) \;=\; \frac{\boldsymbol{d}(t, W)}{\|\boldsymbol{d}(t, W)\|_2} \;\in\; S^{k-1},$$

where S^{k-1} is the unit sphere in $\mathbb{R}^k$.

2.4 Regimes and Regime Boundaries

Definition 2.10 (Regime)**.** A *regime* is a connected, compact subset $\mathcal{R} \subset \Theta$ such that $(t) \in \mathcal{R}$ for a sustained interval of time. Formally, regime $\mathcal{R}$ is *active* at time t if $(t) \in \mathcal{R}$ and there exist $s \leq t \leq e$ with $\min_{t' \in [s,e]} \operatorname{dist}((t'), \mathcal{R}^c) > \epsilon$ for some $\epsilon > 0$ and $e - s \geq L_{\min}$ (a minimum duration parameter).

Definition 2.11 (Regime boundary and transition)**.** The *regime boundary* of $\mathcal{R}$ is the topological boundary $\partial\mathcal{R} \subset \Theta$. A *regime transition at time* t^* occurs when $(t^* - 1) \in \mathcal{R}$ and $(t^*) \notin \mathcal{R}$, i.e., the trajectory exits the regime region. A transition is *smooth* if $\|\boldsymbol{v}(t^*)\|_2$ is bounded, and *discontinuous* (a jump) if $\|\boldsymbol{v}(t^*)\|_2 > \tau$ for a threshold $\tau > 0$.

Definition 2.12 (Jump threshold)**.** For a given significance level $p \in (0, 1)$, the *jump threshold at level* p is the p-th quantile of the empirical distribution of speed:

$$\tau_p \;=\; \inf\bigl\{s : \hat{F}_s(s) \geq p\bigr\},$$

where $\hat{F}_s$ is the empirical CDF of $\{s(t)\}_{t=2}^T$. Observations with $s(t) > \tau_p$ are classified as jump events at level p; observations with $s(t) \leq \tau_p$ constitute the *smooth regime*.

2.5 The Statistical Attractor

Definition 2.13 (Ergodic attractor)**.** For an ergodic process $\{X_t\}$, the *statistical attractor* is

$$\boldsymbol{\theta}^* \;=\; \lim_{T \to \infty} \frac{1}{T} \sum_{t=1}^{T} (t) \;=\; \mathbb{E}_\pi\bigl[(t)\bigr],$$

where π is the stationary measure of the process. In practice, $\boldsymbol{\theta}^*$ is estimated by the sample mean $\hat{\boldsymbol{\theta}}^* = T^{-1} \sum_{t=1}^{T} (t)$.

Definition 2.14 (Attractor distance and stress index)**.** The *attractor distance* at time t is

$$D(t) \;=\; \|(t) - \boldsymbol{\theta}^*\|_2 \,.$$

The *normalised stress index* is

$$\hat{D}(t) \;=\; \frac{D(t)}{\hat{q}_{0.95}(t)},$$

where $\hat{q}_{0.95}(t) = \mathrm{quant}_{0.95}\big\{D(s) : s \le t\big\}$ is the running 95th percentile of the attractor distance, ensuring $\hat{D}(t) \in [0, \infty)$ with $\hat{D}(t) > 1$ indicating unusually high stress.

Definition 2.15 (Rate of stress change)**.** The *stress rate* at time t is

$$\dot{D}(t) \;=\; D(t) - D(t-1).$$

When $\dot{D}(t) > 0$, the system is moving *away* from the attractor (stress is increasing). When $\dot{D}(t) < 0$, the system is *returning* toward the attractor.

2.6 Displacement Persistence

Definition 2.16 (Displacement persistence score)**.** For window size W and $t > W$ and $t + W \le T$, the *displacement persistence score* is

$$P(t, W) \;=\; \langle \hat{\boldsymbol{d}}(t, W),\, \hat{\boldsymbol{d}}(t+W, W) \rangle \;=\; \frac{\boldsymbol{d}(t, W)^\top \boldsymbol{d}(t+W, W)}{\|\boldsymbol{d}(t, W)\|_2 \cdot \|\boldsymbol{d}(t+W, W)\|_2},$$

provided both norms are positive; otherwise $P(t, W) = 0$. This is the cosine similarity between the past displacement direction and the future displacement direction.

Remark 2.2. $P(t, W) \in [-1, 1]$. The sign encodes the dynamical character:

$$P(t,W) > 0 \;\Rightarrow\; \text{momentum}, \qquad P(t,W) = 0 \;\Rightarrow\; \text{random walk}, \qquad P(t,W) < 0 \;\Rightarrow\; \text{mean reversion}.$$

The *directional accuracy* is the fraction of positive scores: $\mathrm{DA}(W) = \mathbb{P}(P(t, W) > 0)$. Mean reversion corresponds to $\mathrm{DA}(W) < 0.5$.

2.7 PCA Reduction and Dimensionality Collapse

In practice we standardise (t) and project to $m \le k$ dimensions via principal component analysis:

$$\hat{(t)} \;=\; \mathbf{U}_m^\top \big((t) - \bar{}\big) \;\in\; \mathbb{R}^m,$$

where $\mathbf{U}_m \in \mathbb{R}^{k \times m}$ contains the top-m eigenvectors of the sample covariance $\hat{\Sigma} = T^{-1} \sum_t ((t) - \bar{})((t) - \bar{})^\top$.

Definition 2.17 (Stress dimensionality)**.** For a local window $[t - L, t + L]$, let $\hat{\lambda}_1(t) \ge \cdots \ge \hat{\lambda}_m(t) \ge 0$ be the eigenvalues of the local sample covariance of $\{\hat{(s)}\}_{s=t-L}^{t+L}$. The

stress dimensionality at time t is

$$\mathrm{sdim}(t) \;=\; \frac{\left(\sum_i \hat{\lambda}_i(t)\right)^2}{\sum_i \hat{\lambda}_i(t)^2} \;\in\; [1, m].$$

This is the effective number of dimensions (participation ratio). Low $\mathrm{sdim}(t)$ indicates that the trajectory has collapsed onto a low-dimensional subspace — the empirical signature of regime stress.

3 Theoretical Results

We now state our main theoretical claims with explicit assumptions and proofs or proof sketches.

3.1 Assumptions

Assumption 3.1 (Bounded parameter space)**.** The statistical parameter manifold Θ is a compact subset of $\mathbb{R}^k$: $\Theta \subset B_M = \{: \|\|_2 \leq M\}$ for some finite $M > 0$.

Assumption 3.2 (Stationarity within regimes)**.** Within any regime $\mathcal{R}$, the process $\{X_t\}$ is weakly stationary: $\mathbb{E}[(t)] = \boldsymbol{\mu}_{\mathcal{R}}$ and $\mathrm{Cov}((t), (t+\ell)) = C_{\mathcal{R}}(\ell)$ for all t and ℓ.

Assumption 3.3 (Ergodicity)**.** The process $\{(t)\}$ is ergodic: time averages converge to ensemble averages, i.e., for any Θ-valued measurable f, $T^{-1}\sum_{t=1}^{T} f((t)) \to \mathbb{E}_\pi[f()]$ almost surely as $T \to \infty$.

Assumption 3.4 (Finite second moments)**.** $\mathbb{E}[\|(t)\|_2^2] < \infty$ for all t.

3.2 Speed as a Regime Transition Signal

Theorem 3.1 (Speed detects transitions)**.** *Let t^* be a regime transition time at which the statistical parameter vector exits regime $\mathcal{R}_1$ and enters $\mathcal{R}_2 \neq \mathcal{R}_1$. Suppose* $\mathrm{dist}(\mathcal{R}_1, \mathcal{R}_2) = \delta > 0$. *Then*

$$s(t^*) \;=\; \|(t^*) - (t^*{-}1)\|_2 \;\geq\; \delta.$$

Moreover, for any non-transition time t with $(t), (t-1) \in \mathrm{int}(\mathcal{R}_1)$, if the within-regime variation is bounded by $\sigma_{\mathcal{R}}$ (i.e., $\|\boldsymbol{v}(t)\|_2 \leq \sigma_{\mathcal{R}}$ a.s.), then $\delta > \sigma_{\mathcal{R}}$ implies $s(t^*) > s(t)$ for all non-transition times t.

Proof. At time t^*, we have $(t^*) \in \mathcal{R}_2$ and $(t^*{-}1) \in \mathcal{R}_1$. By definition of $\mathrm{dist}(\mathcal{R}_1, \mathcal{R}_2)$:

$$s(t^*) = \|(t^*) - (t^*{-}1)\|_2 \geq \inf_{a \in \mathcal{R}_1,\, b \in \mathcal{R}_2} \|b - a\|_2 = \mathrm{dist}(\mathcal{R}_1, \mathcal{R}_2) = \delta.$$

For non-transition times t, $\|\boldsymbol{v}(t)\|_2 \leq \sigma_{\mathcal{R}} < \delta \leq s(t^*)$ by hypothesis, completing the proof. $\square$

Corollary 3.1. *Under continuous observation, the speed $s(t)$ is a necessary indicator of regime transitions: every transition produces $s(t^*) \geq \delta$, while non-transition periods produce $s(t) \leq \sigma_{\mathcal{R}}$. The signal-to-noise ratio of the speed detector is $\delta/\sigma_{\mathcal{R}}$.*

3.3 The Universal Attractor Theorem

Theorem 3.2 (Universal mean reversion in Θ-space)**.** *Let $\{(t)\}$ be a Θ-valued process satisfying Assumptions 3.1–3.4. Then for any displacement window W:*

$$\limsup_{W\to\infty} \mathbb{E}\big[P(t,W)\big] \;\leq\; 0.$$

Furthermore, if the process is not exactly periodic, then for all sufficiently large W:

$$\mathbb{E}\big[P(t,W)\big] \;<\; 0.$$

Proof. Write $\boldsymbol{d}(t,W) = (t) - (t-W)$ and $\boldsymbol{d}(t+W,W) = (t+W) - (t)$. Note that

$$\boldsymbol{d}(t,W) + \boldsymbol{d}(t+W,W) \;=\; (t+W) - (t-W) \;=\; \boldsymbol{d}(t+W,2W).$$

Since Θ is compact (Assumption 3.1):

$$\|\boldsymbol{d}(t+W,2W)\|_2 \;=\; \|(t+W) - (t-W)\|_2 \;\leq\; 2M \quad \text{for all } t, W.$$

Now, by the polarisation identity:

$$\|\boldsymbol{d}(t,W) + \boldsymbol{d}(t+W,W)\|_2^2 = \|\boldsymbol{d}(t,W)\|_2^2 + 2\langle \boldsymbol{d}(t,W),\, \boldsymbol{d}(t+W,W)\rangle + \|\boldsymbol{d}(t+W,W)\|_2^2\,.$$

Therefore:

$$2\langle \boldsymbol{d}(t,W),\, \boldsymbol{d}(t+W,W)\rangle = \|\boldsymbol{d}(t+W,2W)\|_2^2 - \|\boldsymbol{d}(t,W)\|_2^2 - \|\boldsymbol{d}(t+W,W)\|_2^2\,.$$

Taking expectations and using $\mathbb{E}[\|\boldsymbol{d}(t+W,2W)\|_2^2] \leq 4M^2$:

$$2\mathbb{E}\big[\langle \boldsymbol{d}(t,W),\, \boldsymbol{d}(t+W,W)\rangle\big] \leq 4M^2 - \mathbb{E}\big[\|\boldsymbol{d}(t,W)\|_2^2\big] - \mathbb{E}\big[\|\boldsymbol{d}(t+W,W)\|_2^2\big].$$

If $\mathbb{E}[P(t,W)] > 0$, then the expected displacement norm grows as W increases, since $\langle \hat{\boldsymbol{d}}(t,W),\, \hat{\boldsymbol{d}}(t+W,W)\rangle > 0$ means the trajectory tends to continue in the same direction.

Iterating over n non-overlapping windows of width W:

$$\mathbb{E}\big[\|(t+nW)-(t)\|_2\big] \;\to\; \infty \quad \text{as } n \to \infty,$$

contradicting the compactness of Θ (Assumption 3.1). Therefore $\mathbb{E}[P(t,W)] \le 0$ for large W. For strict inequality: if $\mathbb{E}[P(t,W)] = 0$ for all W, the process is a martingale increment process in direction space, which for a compact-valued ergodic process requires exact periodicity. For non-periodic processes, strict inequality holds. □ □

Corollary 3.2 (Existence of a statistical attractor)**.** *Under the conditions of Theorem 3.2, the time-averaged trajectory converges to a point $\boldsymbol{\theta}^* \in \Theta$:*

$$\boldsymbol{\theta}^* \;=\; \lim_{T\to\infty} \frac{1}{T} \sum_{t=1}^{T} (t)$$

and the expected future displacement is directed toward $\boldsymbol{\theta}^*$:

$$\mathbb{E}\big[(t+W)-(t) \,\big|\, (t)\big] \;=\; f_W\big((t)\big)$$

where $\langle f_W(), -\boldsymbol{\theta}^*\rangle \le 0$ for all $\in \Theta$ and large W (the drift points toward the attractor).

3.4 Ornstein-Uhlenbeck Representation

The empirical mean-reversion is consistent with the following continuous-time representation.

Proposition 3.1 (OU representation of (t))**.** *Under Assumptions 3.1–3.4, if the trajectory $\{(t)\}$ is generated by a diffusion process on Θ, then the process admits the Ornstein-Uhlenbeck representation:*

$$d(t) \;=\; \kappa\big(\boldsymbol{\theta}^* - (t)\big)\,dt \;+\; \boldsymbol{\Sigma}^{1/2}\,d\mathbf{W}(t),$$

where $\kappa > 0$ is the mean-reversion speed, *$\boldsymbol{\Sigma} \in \mathbb{R}^{k\times k}$ is a positive semi-definite diffusion matrix, and $\mathbf{W}(t)$ is a standard k-dimensional Brownian motion. Under this model, the displacement persistence satisfies:*

$$\mathbb{E}\big[P(t,W)\big] \;=\; -\big(1-e^{-2\kappa W}\big) \;\xrightarrow{W\to\infty}\; -1.$$

Thus $\mathrm{DA}(W) < 0.5$ *for all $W > 0$, with the gap* $0.5 - \mathrm{DA}(W)$ *increasing monotonically with W.*

Proof. Under the OU model, the displacement $\boldsymbol{d}(t,W)$ has conditional mean

$$\mathbb{E}\big[\boldsymbol{d}(t,W) \mid (t-W)\big] = \big(e^{-\kappa W}-1\big)\big((t-W)-\boldsymbol{\theta}^*\big),$$

which points from $(t - W)$ toward $\boldsymbol{\theta}^*$, i.e., opposite to $(t - W) - \boldsymbol{\theta}^*$. Similarly for $\boldsymbol{d}(t + W, W)$ with respect to (t). Computing the expected cosine similarity under the OU model and taking leading-order terms gives the stated expression. Full derivation follows from Uhlenbeck & Ornstein [1930]. □ □

3.5 Attractor Escape and Crisis Regimes

Definition 3.1 (Escape phase)**.** The system is in an *escape phase* at time t if all three conditions hold simultaneously:

(i) $D(t) > q_{0.95}$ (far from attractor),

(ii) $\dot{D}(t) > 0$ (distance increasing),

(iii) $\mathbb{E}[P(t, W)] > 0$ locally (momentum dominates).

Proposition 3.2 (Escape implies regime transition)**.** *If the system enters an escape phase at time t_0 and remains in it for $L \geq L_{\min}$ consecutive periods, then either a regime transition occurs at some $t^* \in [t_0, t_0 + L]$, or the system has entered a new attractor basin $\mathcal{R}'$ with attractor $\boldsymbol{\theta}^{*\prime} \neq \boldsymbol{\theta}^*$.*

Proof. By Theorem 3.2, sustained positive persistence $\mathbb{E}[P(t, W)] > 0$ over L periods implies $\mathbb{E}[\|(t_0 + L) - (t_0)\|_2] \geq c\sqrt{L}$ for some $c > 0$. For large enough L, this displacement exceeds the diameter of any regime region $\mathcal{R}$, so the trajectory must have exited the current regime. □ □

3.6 The Dimensionality Collapse Theorem

Theorem 3.3 (Dimensionality collapse under stress)**.** *Let $\hat{\Sigma}(t)$ be the local covariance of the trajectory over a window $[t - L, t + L]$ and let $\lambda_1(t) \geq \cdots \geq \lambda_m(t)$ be its eigenvalues. Under an escape phase (Definition 3.1), the trajectory concentrates on a low-dimensional subspace:*

$$\mathrm{sdim}(t) \;\to\; 1 \quad \textit{as the escape accelerates.}$$

Specifically, if the escape is along a principal direction $\boldsymbol{u}_1 \in S^{k-1}$, then $\lambda_1(t)/\sum_i \lambda_i(t) \to 1$ as the escape dominates the within-window variation.

Proof. During an escape, $\boldsymbol{v}(t) \approx c(t)\boldsymbol{u}_1$ for some scalar $c(t) > 0$ (the trajectory moves predominantly along $\boldsymbol{u}_1$). The local covariance is approximately $\hat{\Sigma}(t) \approx \sigma_c^2 \boldsymbol{u}_1 \boldsymbol{u}_1^\top + \epsilon \mathbf{I}$ where $\sigma_c^2 = \mathrm{Var}(c(t))$ and ϵ is residual noise. For $\sigma_c^2 \gg \epsilon$: $\lambda_1(t) \approx \sigma_c^2 + \epsilon \approx \sigma_c^2$, $\lambda_i(t) \approx \epsilon$ for $i \geq 2$, so $\mathrm{sdim}(t) \approx 1$. □ □

3.7 Attractor Migration

Definition 3.2 (Attractor migration)**.** The attractor $\boldsymbol{\theta}^*$ *migrates* if the stationary measure π_t of the process is itself time-varying: $\boldsymbol{\theta}^*(t) = \mathbb{E}_{\pi_t}[(t)]$ satisfies $\|\boldsymbol{\theta}^*(t) - \boldsymbol{\theta}^*(s)\|_2 > 0$ for $t \neq s$. The *migration rate* is $\|d\boldsymbol{\theta}^*/dt\|_2$.

Remark 3.1. Attractor migration is empirically identifiable as a systematic trend in the attractor distance $D(t)$ that is not reversible on any observed timescale. It represents a structural change in the long-run statistical character of the system and requires a time-varying extension of Definition 2.13.

4 The Data-Trust Score $\alpha(t)$

4.1 Motivation

Any statistical model $\mathcal{M}$ trained on historical data implicitly assumes that the current distribution p_t is well-approximated by the training distribution p_{train}. This assumption fails precisely when (t) is far from the historical attractor $\boldsymbol{\theta}^*$. We formalise this through a continuous trust score.

4.2 Formal Definition

Definition 4.1 (Data-trust score)**.** The *data-trust score* at time t is

$$\alpha(t) \;=\; \sigma\big(-\beta_1 \hat{D}(t) - \beta_2 \dot{D}(t)\big) \;\in\; (0,1), \tag{2}$$

where $\sigma(x) = (1+e^{-x})^{-1}$ is the logistic function, $\hat{D}(t)$ is the normalised attractor distance (Definition 2.14), $\dot{D}(t)$ is the stress rate (Definition 2.15), and $\beta_1, \beta_2 > 0$ are calibration parameters.

Proposition 4.1 (Properties of $\alpha(t)$)**.** *The trust score $\alpha(t)$ defined in* (2) *satisfies:*

(i) $\alpha(t) \to 1$ *when* $\hat{D}(t) \to 0$ *and* $\dot{D}(t) \to 0$ *(near attractor, not moving): full trust.*

(ii) $\alpha(t) \to 0$ *when* $\hat{D}(t) \to \infty$ *or* $\dot{D}(t) \to \infty$ *(far from attractor, escaping): minimal trust.*

(iii) $\alpha(t)$ *is monotone decreasing in both* $\hat{D}(t)$ *and* $\dot{D}(t)$*.*

(iv) $\partial\alpha/\partial\hat{D} = -\beta_1\alpha(1-\alpha)$*: the sensitivity to attractor distance is highest at intermediate stress levels.*

4.3 Blended Prediction Under Uncertainty

Let $\hat{Y}_{\text{data}}(t+h)$ be a data-driven forecast and $\hat{Y}_{\text{prior}}(t+h)$ be a theory-based structural forecast.

Definition 4.2 (Trust-weighted prediction)**.** The *trust-weighted prediction* at horizon h is

$$\hat{Y}(t+h) \;=\; \alpha(t)\cdot\hat{Y}_{\text{data}}(t+h) \;+\; \big(1-\alpha(t)\big)\cdot\hat{Y}_{\text{prior}}(t+h).$$

Remark 4.1. This formalises the intuition that *during a crisis, historical statistical relationships break down and structural priors should dominate.* When $\alpha(t)\approx 1$ (near attractor), the data model dominates. When $\alpha(t)\approx 0$ (escape phase), the structural prior dominates.

4.4 Calibration

The parameters β_1 and β_2 can be calibrated by minimising the out-of-sample prediction error of the blended model over a training period:

$$(\hat{\beta}_1,\hat{\beta}_2) \;=\; {}_{\beta_1,\beta_2>0}\sum_{t\in\mathcal{T}_{\text{val}}} \ell\big(Y(t),\hat{Y}(t)\big),$$

where ℓ is a suitable loss function (squared error, log-loss, etc.) and $\mathcal{T}_{\text{val}}$ is a validation period.

5 Empirical Methodology

5.1 Datasets

Table 1: Datasets used in cross-domain validation.

Domain	Source	Resolution	Span	Window W_θ
S&P 500 returns	Yahoo Finance	Daily	1994–2024	60 days
Sunspot activity	SILSO / Royal Obs.	Weekly	1960–2024	52 weeks
US COVID-19 cases	Our World in Data	Daily	2020–2023	21 days
Global temp. anomaly	NASA GISS	Monthly	1880–2024	24 months

5.2 Statistical Tests

For each domain we perform four tests aligned with the theoretical claims:

Test A (Speed detection). For each known event time t^*, compute the percentile rank of $s(t^*)$ in the empirical distribution of $\{s(t)\}$. Values above the 90th percentile constitute a detection.

Test B (Displacement persistence). Compute $P(t, W)$ for all valid t and test $H_0 : \mathbb{E}[P] = 0$ vs $H_1 : \mathbb{E}[P] < 0$ using a one-sample t-test.

Test C (Attractor distance at events). Test whether $D(t)$ is significantly elevated during known stress periods using a two-sample t-test against baseline periods.

Test D (Directional accuracy). Compute $\mathrm{DA}(W) = \hat{\mathbb{P}}(P(t, W) > 0)$. Values below 0.5 confirm mean reversion.

6 Empirical Results

6.1 Cross-Domain Displacement Persistence

Table 2: Displacement persistence test across four domains. All reject $H_0 : \mathbb{E}[P] = 0$ in favour of mean reversion at $p < 0.001$.

Domain	W	n	$\hat{\mathbb{E}}[P]$	$\mathrm{DA}(W)$	t-stat	p-value
S&P 500 (finance)	20d	7,702	-0.039	0.473	-7.21	$< 10^{-3}$
Sunspot activity	26wk	3,054	-0.046	0.458	-19.4	$< 10^{-3}$
COVID-19 (US)	14d	694	-0.137	0.399	-22.8	$< 10^{-3}$
Global temperature	12mo	1,608	-0.151	0.355	-46.3	$< 10^{-3}$

The universal negative sign of $\hat{\mathbb{E}}[P]$ across all four domains, spanning finance, solar physics, epidemiology, and climatology, constitutes empirical confirmation of Theorem 3.2. The gradient in strength (finance weakest, temperature strongest) is consistent with the theoretical expectation: financial data is most contaminated by exogenous shocks that temporarily override the restoring force, while temperature and epidemic dynamics are more purely governed by bounded internal dynamics.

6.2 Escape Phase Detection: Bear 2020 Q1

Within the financial domain, the Covid crash (Bear 2020 Q1) is the clearest empirical example of an escape phase (Definition 3.1). During this six-week period, the directional accuracy is 0.833 — strongly positive persistence, consistent with Proposition 3.2. The attractor distance $D(t)$ reaches its maximum in the 30-year record during this period. This constitutes evidence that the three-condition escape phase criterion correctly identifies qualitatively different dynamics from the mean-reverting baseline.

6.3 Dimensionality Collapse

Under regime stress in the S&P 500 data, the proportion of total variance explained by PC1 rises from approximately 40% in normal periods to 85% during the GFC and Covid

crash, consistent with Theorem 3.3. The first principal component under stress loads primarily on σ_t^2, κ_t, and $\sigma_{\downarrow,t}^2$, indicating that the dominant direction of instability is the volatility-kurtosis axis.

6.4 Attractor Migration in Climate Data

The global temperature dataset shows a systematic upward trend in $D(t)$ from approximately 1970 onward, consistent with Definition 3.2. The local mean-reverting El Niño oscillations remain present (negative local persistence at 12-month scale) but are superimposed on a migrating attractor whose rate is itself non-stationary. This is the geometric representation of anthropogenic climate change in the SMA framework.

7 Discussion

7.1 What is New

The SMA framework makes three contributions that, to our knowledge, do not exist in the literature as a unified whole: (i) the representation of regime state as a continuous trajectory on a statistical parameter manifold rather than as a discrete label; (ii) the proof and cross-domain empirical demonstration that displacement persistence is universally negative in bounded systems; and (iii) the derivation of a data-trust score $\alpha(t)$ that translates the geometric state of the manifold into an actionable discount factor for downstream prediction.

7.2 Relation to Existing Work

The negative displacement persistence is consistent with OU dynamics (Proposition 3.1) but derived here from first principles (compactness of Θ) without parametric assumptions. The early warning signals literature [Scheffer et al., 2009, Dakos et al., 2012] identifies rising variance and autocorrelation as precursors to critical transitions; in our framework these appear as speed increases and dimensionality collapse, respectively. BOCPD [Adams & MacKay, 2007] provides a Bayesian posterior over changepoint times; our framework provides a continuous real-time geometric characterisation of the same phenomenon.

7.3 Limitations

The attractor estimate * is biased under attractor migration (Section 6). The displacement window W must be chosen relative to the natural timescale of the domain. The trust score calibration parameters β_1, β_2 require domain-specific tuning. The framework in its

current form is developed for scalar processes; multivariate extension is direct but not yet validated.

8 Conclusion

We have introduced the Statistical Manifold Attractor (SMA) Framework, a geometric approach to non-stationarity in which the statistical state of any time series is represented as a trajectory (t) on a parameter manifold Θ. We proved that displacement persistence is universally non-positive for any bounded dynamical system (Theorem 3.2), validated this across four independent domains (finance, solar physics, epidemiology, climatology), and derived a data-trust score $\alpha(t)$ that translates manifold geometry into a principled discount factor for any downstream model.

The central claim is simple and general:

Any bounded dynamical system has a statistical attractor. Regime stress is departure from that attractor.
Regime shift is escape to a new attractor basin.
The data-trust score $\alpha(t)$ measures how far and how fast the system is departing from what it knows.

References

Adams, R. P., & MacKay, D. J. C. (2007). Bayesian online changepoint detection. *arXiv:0710.3742*.

Amari, S. (1985). *Differential-Geometrical Methods in Statistics*. Springer, New York.

Dakos, V., Carpenter, S. R., Brock, W. A., Ellison, A. M., Guttal, V., Ives, A. R., Kéfi, S., Livina, V., Seekell, D. A., van Nes, E. H., & Scheffer, M. (2012). Methods for detecting early warnings of critical transitions in time series illustrated using simulated ecological data. *PLOS ONE*, 7(7), e41010.

Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357–384.

Killick, R., Fearnhead, P., & Eckley, I. A. (2012). Optimal detection of changepoints with a linear computational cost. *Journal of the American Statistical Association*, 107(500), 1590–1598.

Poterba, J. M., & Summers, L. H. (1988). Mean reversion in stock prices: Evidence and implications. *Journal of Financial Economics*, 22(1), 27–59.

Runge, V., Hocking, T. D., Romano, G., Installe, A., Fearnhead, P., & Rigaill, G. (2020). gfpop: An R package for univariate graph-constrained change-point detection using dynamic programming. *arXiv:2002.03646*.

Scheffer, M., Bascompte, J., Brock, W. A., Brovkin, V., Carpenter, S. R., Dakos, V., Held, H., van Nes, E. H., Rietkerk, M., & Sugihara, G. (2009). Early-warning signals for critical transitions. *Nature*, 461, 53–59.

Uhlenbeck, G. E., & Ornstein, L. S. (1930). On the theory of the Brownian motion. *Physical Review*, 36(5), 823–841.

Vovk, V., Gammerman, A., & Shafer, G. (2005). *Algorithmic Learning in a Random World*. Springer, New York.